

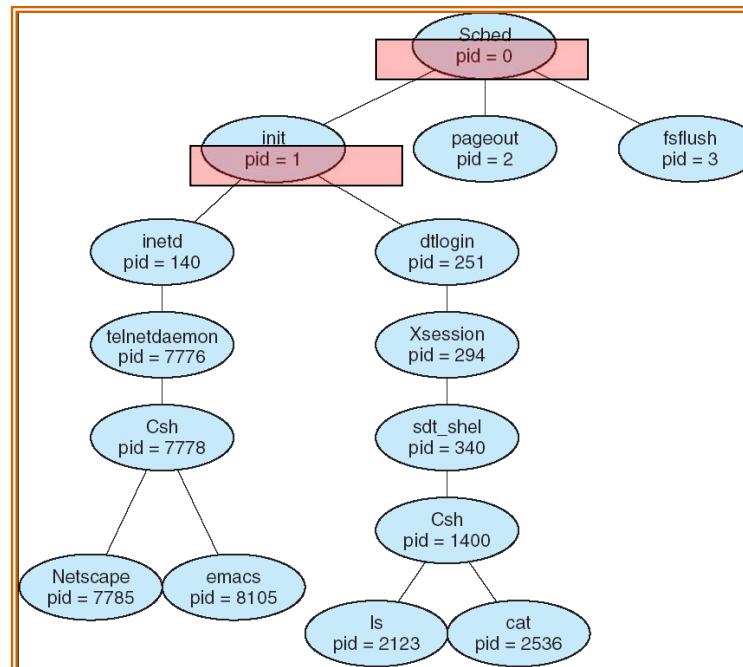
## Operations on processes

---

- Process creation
- Process termination

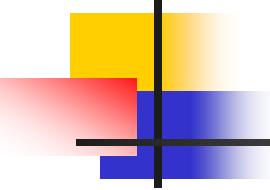
- Parent process create children processes, which, in turn create other processes, forming a tree of processes

**PID**



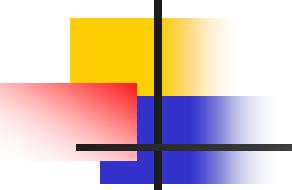
**Solaris**





## ■ What are the issues?

- Resource sharing options
  - 1. Parent and child share all resources
  - 2. Child shares subset of parent's resources
  - 3. Parent and child share no resources
- Execution options
  - 1. Parent and child execute concurrently
  - 2. Parent waits until child terminates
- Address space
  - 1. Child duplicates parent's address space
  - 2. Child has a program loaded into it

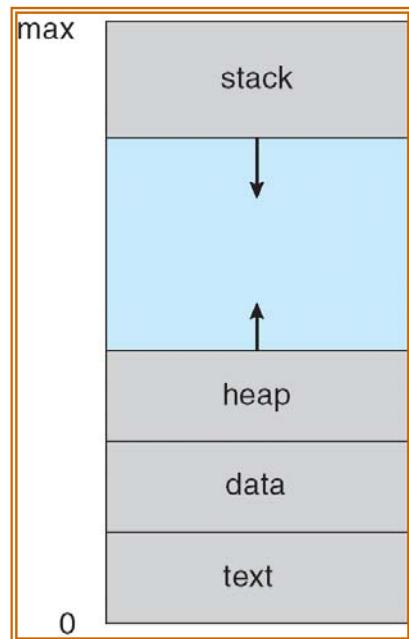


## ■ So, in case of UNIX or Linux

- **fork system call**

- Resource sharing
  - Parent and children share **files**
  - But CPU time or memory are not shared
- Execution
  - Parent and children execute concurrently
- Address space
  - Duplicate but separate address spaces
  - At first, each process has the same user stack, the same local variables, the same heap and the same program
- Call once, return twice
  - Fork is called by parent, but it returns twice: once to the parent (PID of child) and to the child (0)

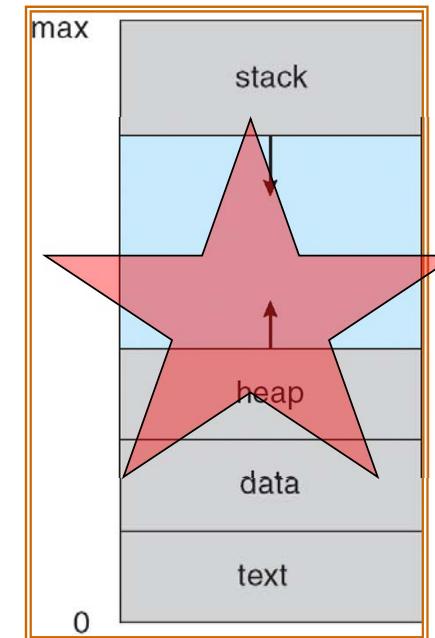
## Parent process

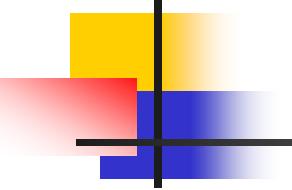


**At first, duplicate  
the address space**

**But, typically  
overwrite the new  
program onto it**

## Child process



- 
- Every process has the same program !

**exec** system call is used to load a new program

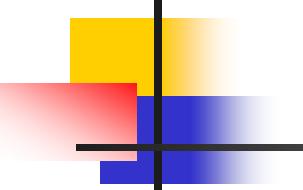
exec system call loads a program and initializes the program

- Reaping child processes
  - Wait system call

A process waits for its children to terminate

What does “zombie process” mean ?  
A terminated process that has not yet been reaped





```
int main()
{
    pid_t pid;
    /* fork another process */
    pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execvp("/bin/ls", "ls", NULL);
    }
    else /* parent process */
    {
        /* parent will wait for the child to complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

## Call once Return twice



```
int main()
{
    pid_t pid;
    /* fork another process */
    → pid = fork();
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
    else if (pid == 0) { /* child process */
        execvp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf ("Child Complete");
        exit(0);
    }
}
```

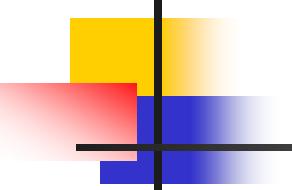
**parent**

**PC** →

**/bin/ls program**

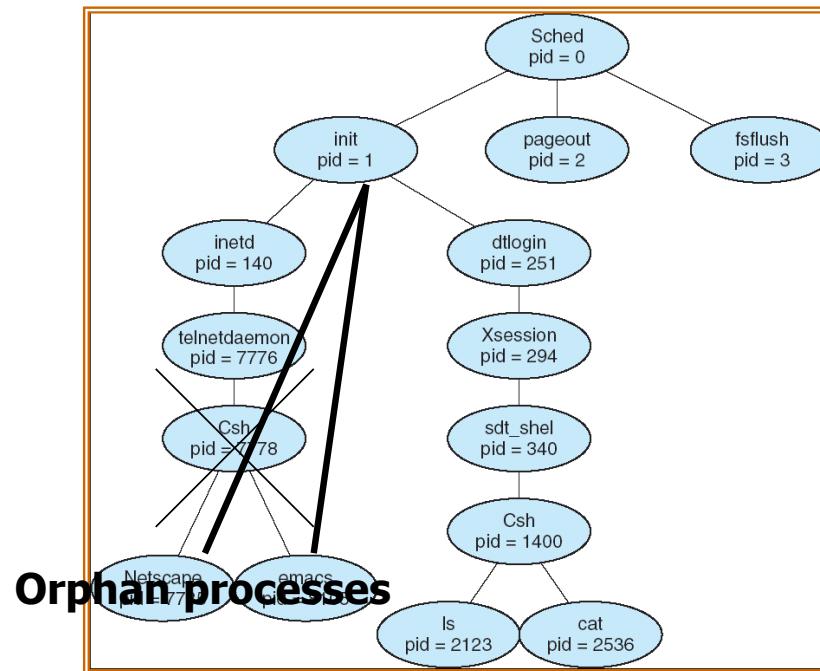
**child**

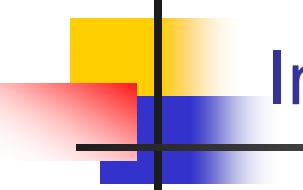




## ■ Process termination

- Process executes last statement and asks the operating system to delete it (**exit**)
  - Output data from child to parent (via **wait**)
  - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (**abort**)
  - Child has exceeded allocated resources
  - Task assigned to child is no longer required
  - If parent is exiting
    - Some operating system do not allow child to continue if its parent terminates
      - All children terminated - *cascading termination*
    - **Orphan process**
      - **What does OS do?**





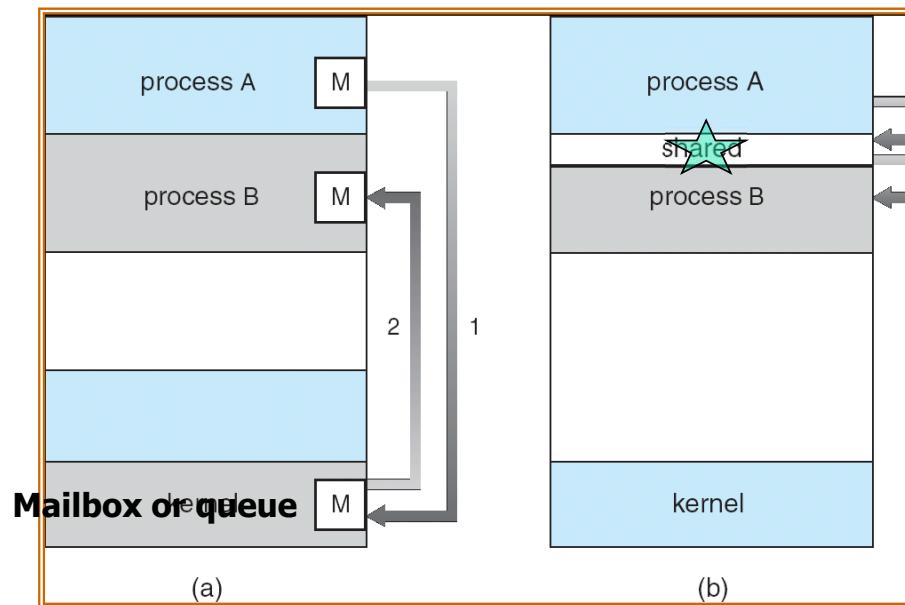
# Inter-process communication (IPC)

- **Independent** process cannot affect or be affected by the execution of another process
- **Cooperating** process can affect or be affected by the execution of another process
- Advantages of process cooperation
  - Information sharing
    - e.g. shared file
  - Computation speed-up
    - Parallel processes <if the computer has multiple processing elements>
  - Modularity
  - Convenience
    - e.g. editing printing compiling at once



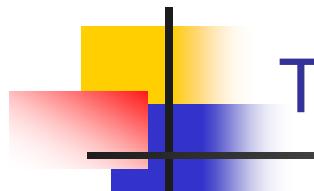
## ■ IPC

- 1. Message passing
- 2. Shared memory



**Message passing**

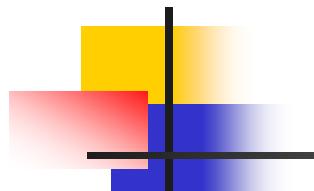
**Shared memory**



## Typical examples of IPC

- Producer-Consumer problem
  - Paradigm for cooperating processes, a *producer* process produces information that is consumed by a *consumer* process
    - *unbounded-buffer* places no practical limit on the size of the buffer
    - *bounded-buffer* assumes that there is a fixed buffer size
  - Must be synchronized





- Producer-consumer problem using *a programmed shared-memory*

- Bounded buffer

- Shared data

```
#define BUFFER_SIZE 10
```

```
typedef struct {
```

```
    ...
```

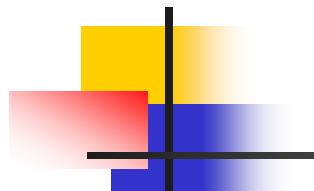
```
} item;
```

```
item buffer[BUFFER_SIZE]; // A circular buffer
```

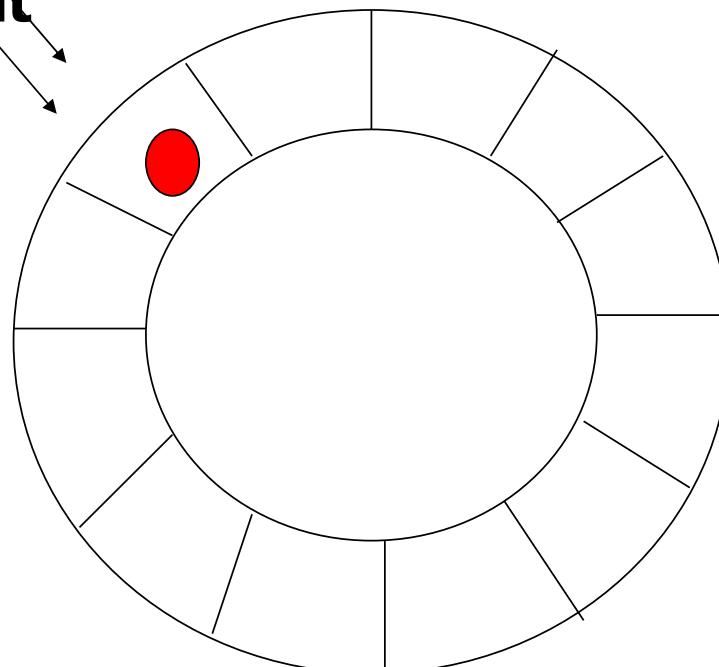
```
int in = 0; // Free position in the buffer
```

```
int out = 0; // The first full position in the buffer
```





**in**  
**out**



**Producer**

**Producing  
item**

**Consumer**

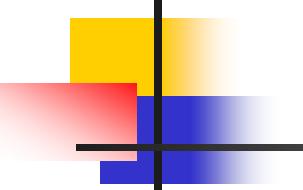
**Waiting  
Spin lock**

**Consuming  
item**

**Empty case:  $\text{in}=\text{out}$**

**Shared memory, bounded buffer**





```
while (true) {
    /* Produce an item */
    while (((in + 1) % BUFFER SIZE) == out)
        ; /* do nothing - no free buffers */
    buffer[in] = item;
    in = (in + 1) % BUFFER SIZE;
}
```

Producer

Consumer

```
while (true) {
    while (in == out)
        ; // do nothing -- nothing to consume

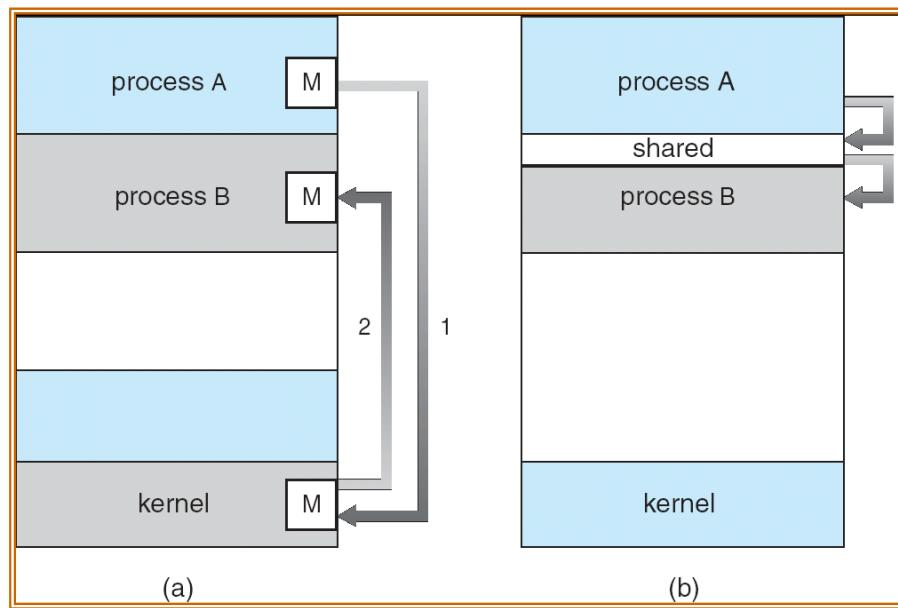
    // remove an item from the buffer
    item = buffer[out];
    out = (out + 1) % BUFFER SIZE;
    return item;
}
```

Synchronization !



## ■ IPC

- 1. Message passing
- 2. Shared memory

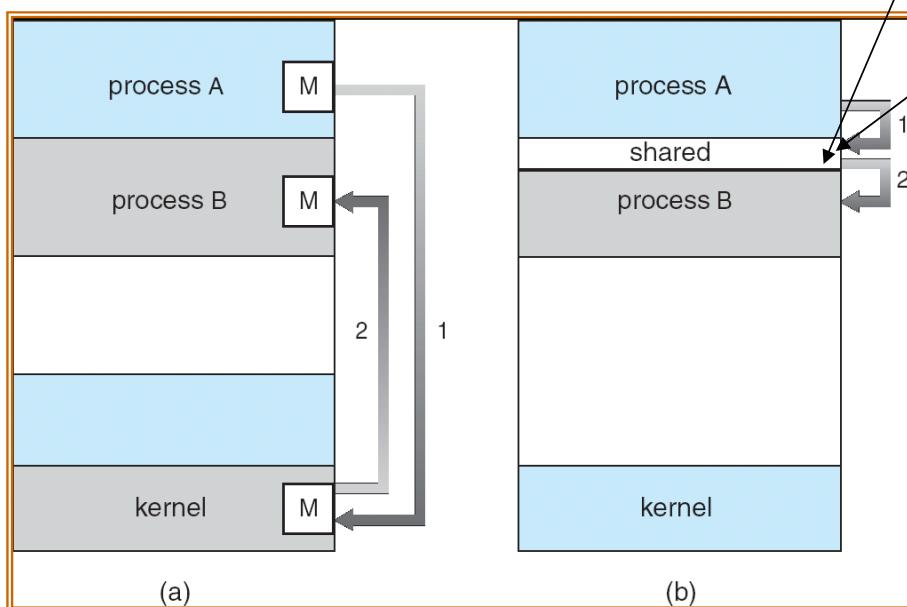


**Message passing**

**Shared memory**

## ■ Shared memory

**Address space of the process creating  
the shared memory region**



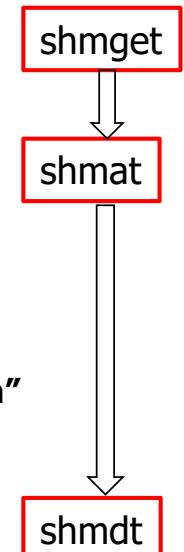
**1. Process B creates shared region**

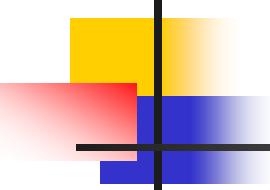
**2. Process A attaches the region to  
its own address space**

**3. Process B writes "Shared region"  
in the shared region**

**4. Process A can read "Shared region"  
in the shared region**

**5. Detach a shared region**





## ■ Message-passing

- Processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
  - **send**(message) – message size fixed or variable
  - **receive**(message)
- If  $P$  and  $Q$  wish to communicate, they need to:
  - 1. establish a *communication link* between them
  - 2. exchange messages via send/receive
- Implementation of communication link
  - physical (e.g., shared memory, hardware bus)
  - logical (e.g., logical properties)



## ■ Establishing links

- 1. Direct communication

- Processes must **name each other explicitly**:
    - **send** ( $P, \text{ message}$ ) – send a message to process P
    - **receive**( $Q, \text{ message}$ ) – receive a message from process Q



**P**

**OS**



**Q**

- 2. Indirect communication

- Messages are **directed and received from mailboxes** (also referred to as ports)
    - Each mailbox has a unique id
    - Processes can communicate only if they share a mailbox



**P**

**OS**



**OS**



**Q**



20