

Operating Systems

(Threads & Concurrency)

Chapter 4

These lecture materials are modified from the source lecture notes
written by A. Silberschatz, P. Galvin and G. Gagne.

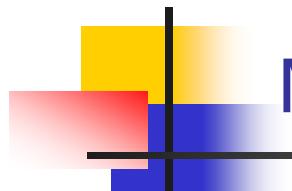
Spring, 2020





Outline

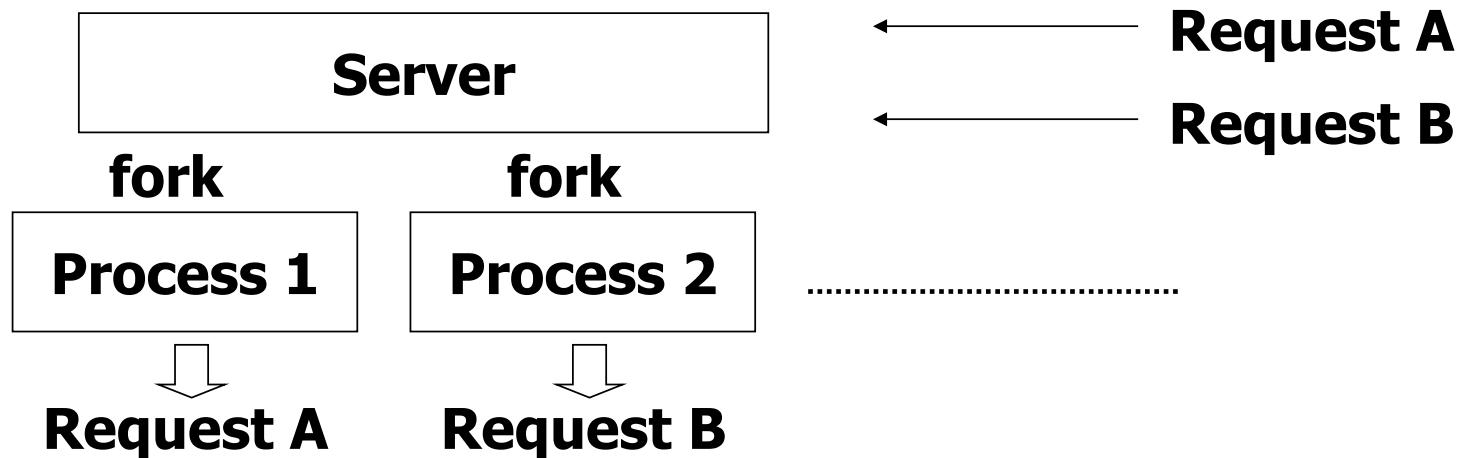
- Motivation
- Overview
- Multicore programming
- Multithreading models
- Thread libraries
- Threading issues
- Thread programming API

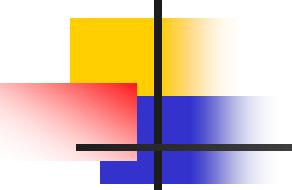


Motivation

Many programs must perform several tasks that do not need to be serialized.

Consider a web server;





**But process creation is
resource-intensive and time-consuming**

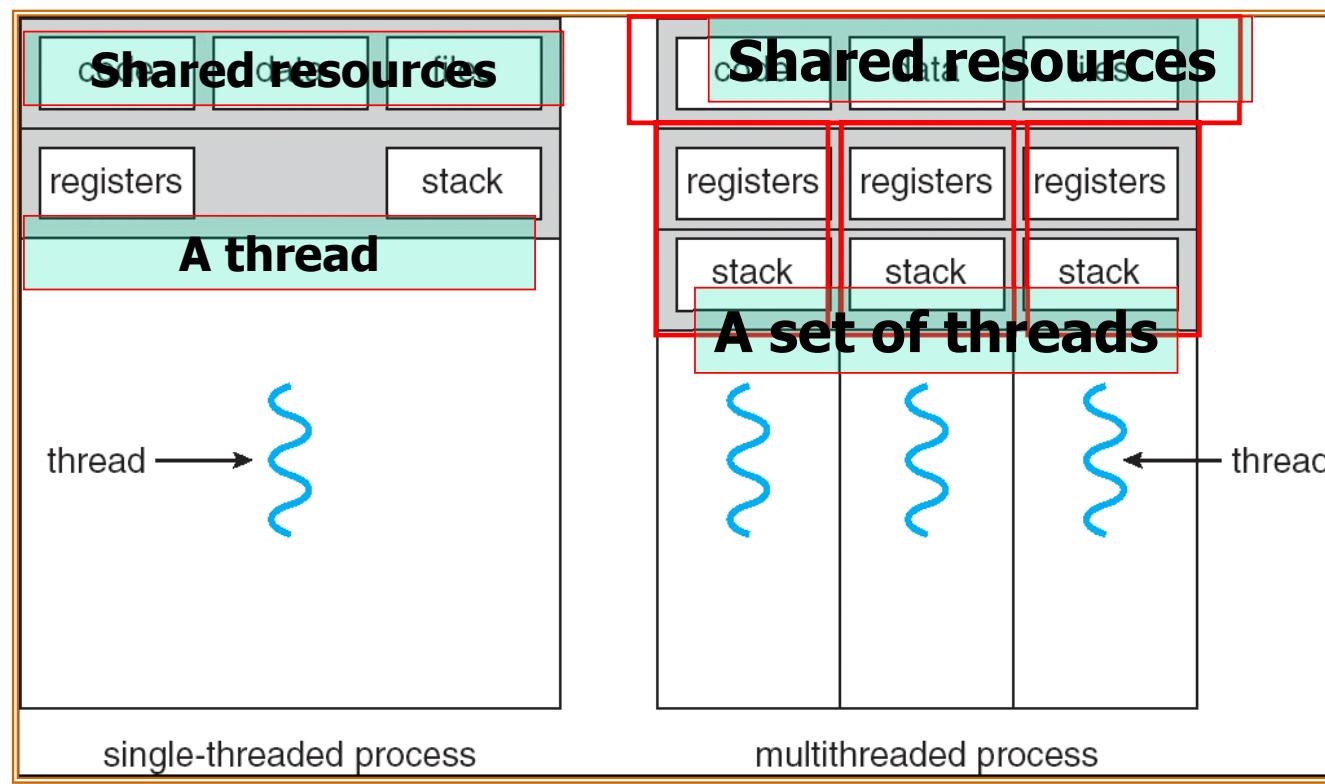
- 1. Allocate a new PID, PCB and fill them in**
 - 2. Duplicate the address space**
 - 3. Make the child runnable**
 - 4. ...**

In addition, an IPC mechanism is needed

**Therefore, a multithreading mechanism that
reduces such overhead is required**

Overview

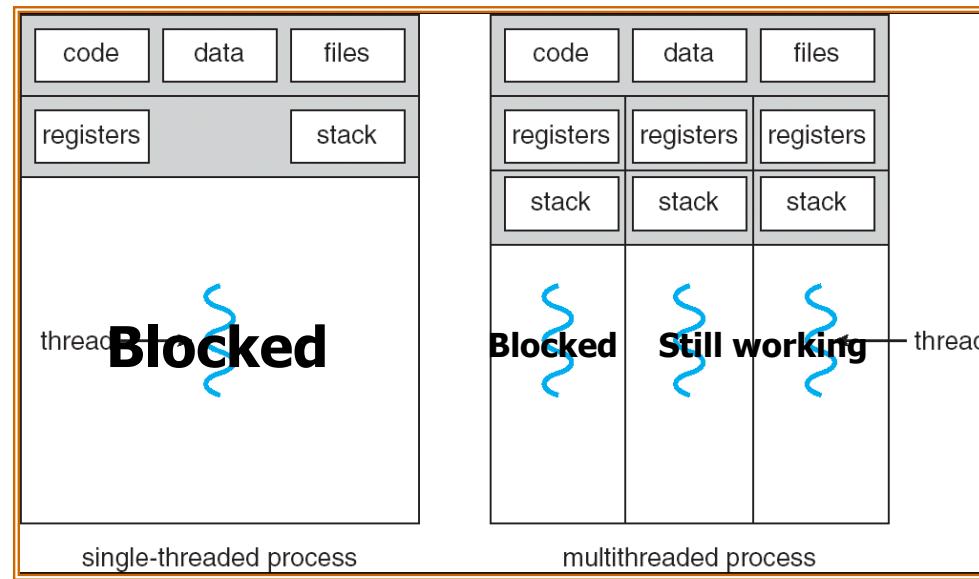
- Single vs. multi-threaded processes



■ Benefits

• 1. Responsiveness

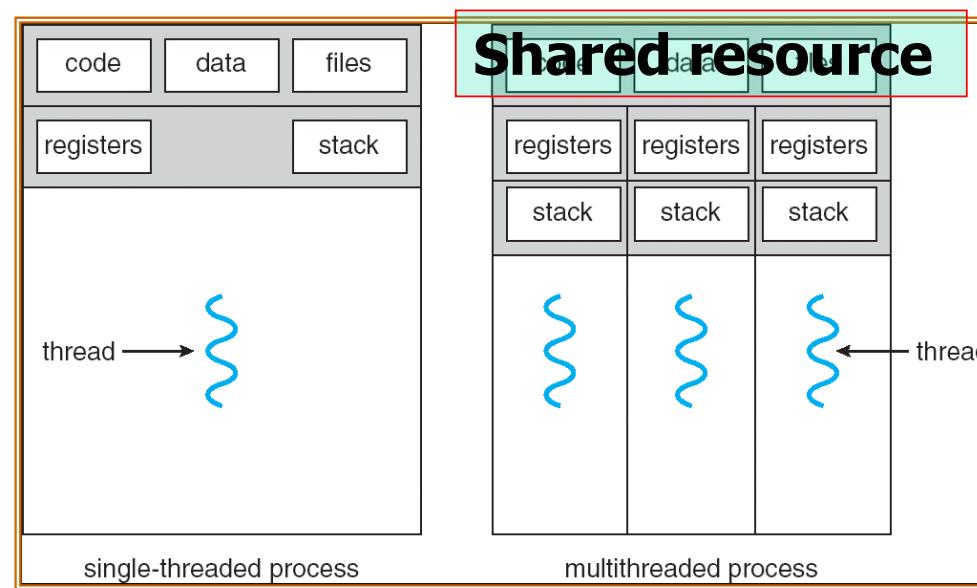
- Multithreading allows a program to continue running if part of it is blocked or is performing a lengthy operation

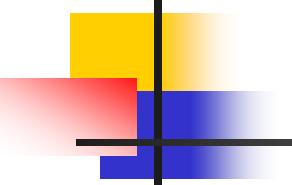


■ Benefits

• 2. Resource Sharing

- Threads share binary code, data, resource of the process





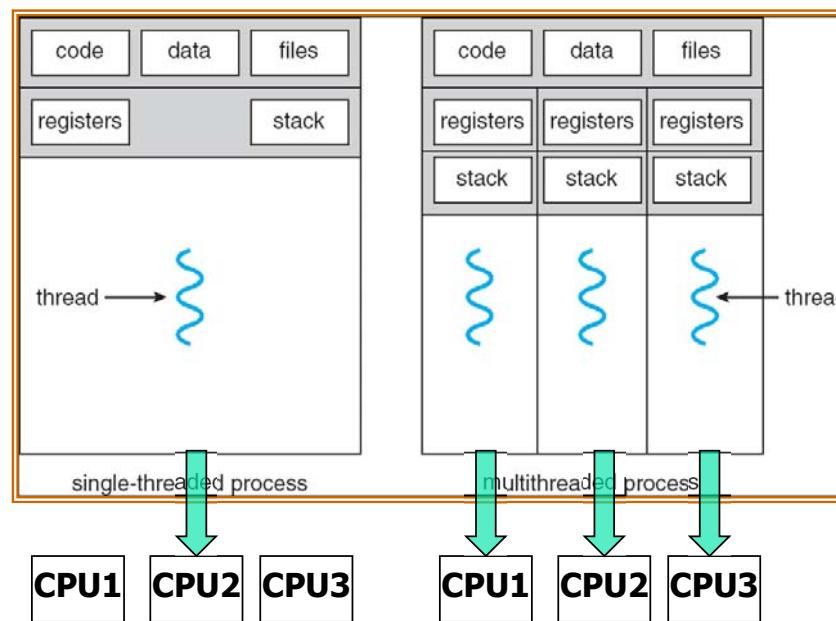
■ Benefits

- 3. Economy

- Allocating memory and resources for **process creation** is costly
- But it is more economical to create or context-switch **threads**
- For example, in Solaris
 - Process creation: 30 times slower than thread creation
 - Process context switch: 5 times slower than thread switch

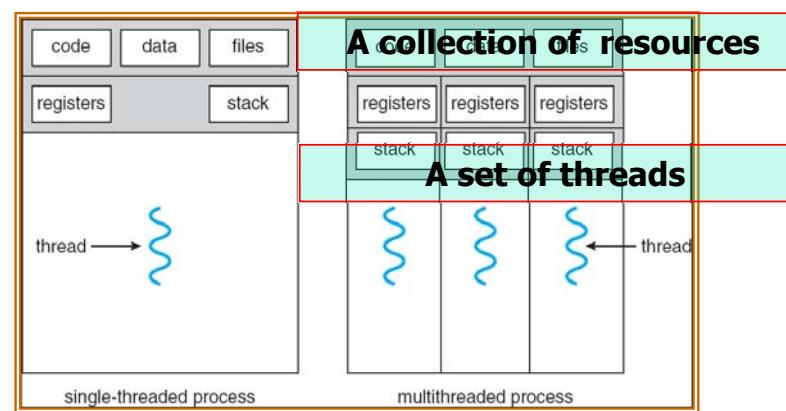
■ Benefits

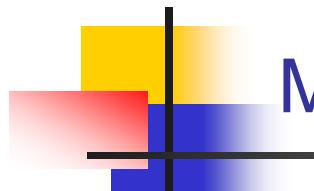
- 4. Utilization of MP (multi-processor) Architectures
 - Threads may be running in parallel on different processors



■ Fundamental concepts

- A process is a compound entity that can be divided into:
 - **A set of threads**
 - A dynamic object that represents a control point in the process
 - **A collection of resources**
 - An address space, open files and so on





Multicore programming

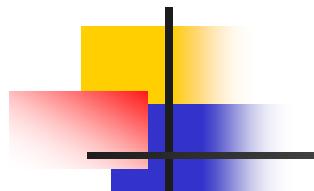
- Types of parallelism
 - **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
 - **Task parallelism** – distributing threads across cores, each thread performing unique operation

```
for(i=1;i<10000000;i++)  
    a+=i;
```

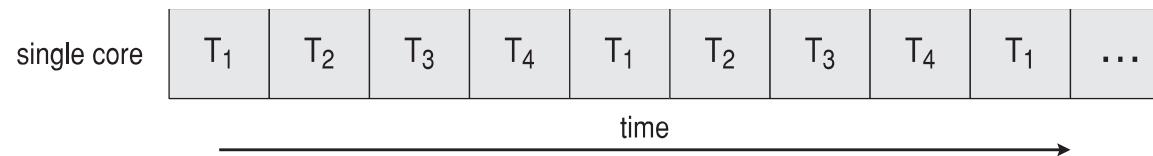


```
for(i=1;i<5000000;i++)  
    a+=i;
```

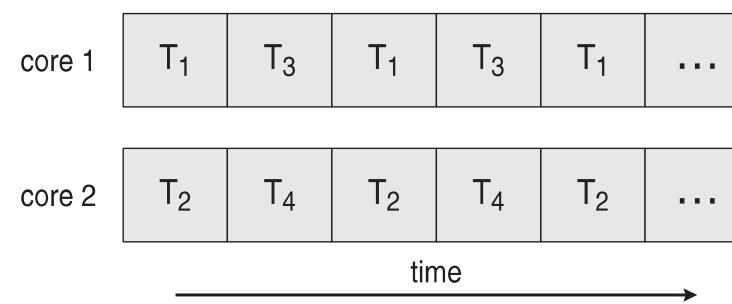
```
for(i=5000000;i<10000000;i++)  
    a+=i;
```

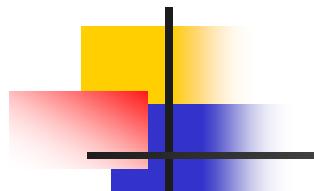


Concurrent execution on single-core system: **(동시성, 병행성)**



Parallelism on a multi-core system: **(병렬성)**



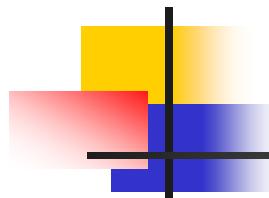


■ Why multicore programming is difficult ?

- 1. Identifying task
- 2. Balance
- 3. Data splitting
- 4. Data dependency
- 5. Testing and debugging

New approach to software design is needed !

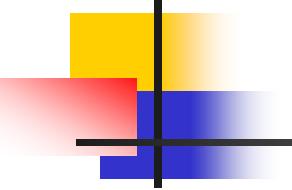




Multi-Threading Model

- There are three types of threads;
 - Kernel thread
 - User thread



- 
- Kernel thread
 - It is created and destroyed as needed internally by the kernel
 - It is managed by the operating system
 - User thread
 - It is supported above the kernel and managed by the **thread library**
 - It does not involve kernel



ka,kb : kernel threads

UA,UB,UC,UD,UE : User threads

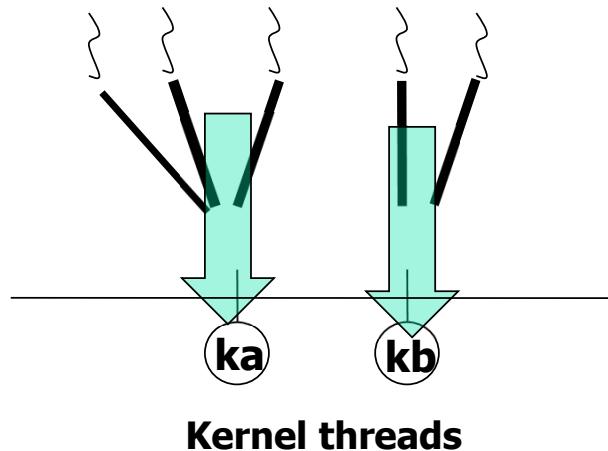
A scheduler is not aware of the existence of user threads

If a kernel thread “ka” is blocked, then user threads UA,UB and UC are blocked as well

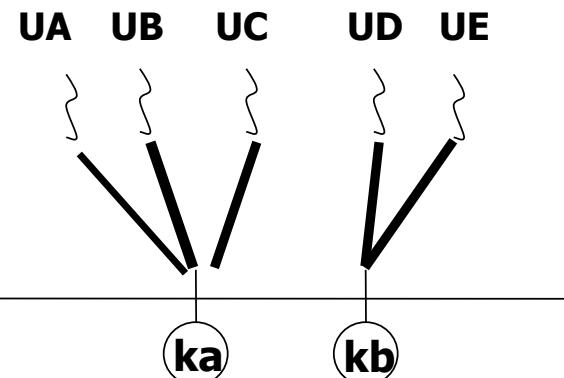
One kernel thread can be allocated to one processor

User threads

UA UB UC UD UE

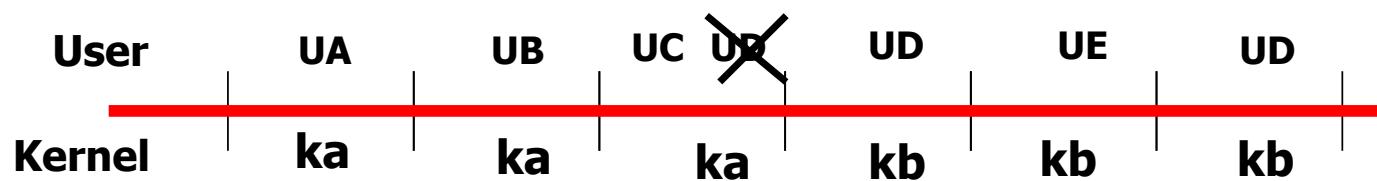


User threads



A scheduler is not aware of the existence of user threads

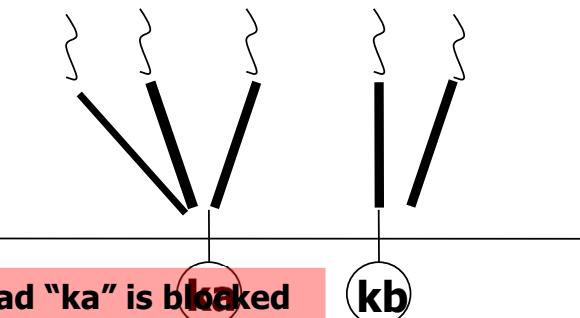
Kernel threads



User threads

User threads will be blocked UA UB UC

UD UE



If a kernel thread is blocked, then its associated user threads are blocked as well.

A thread "ka" is blocked

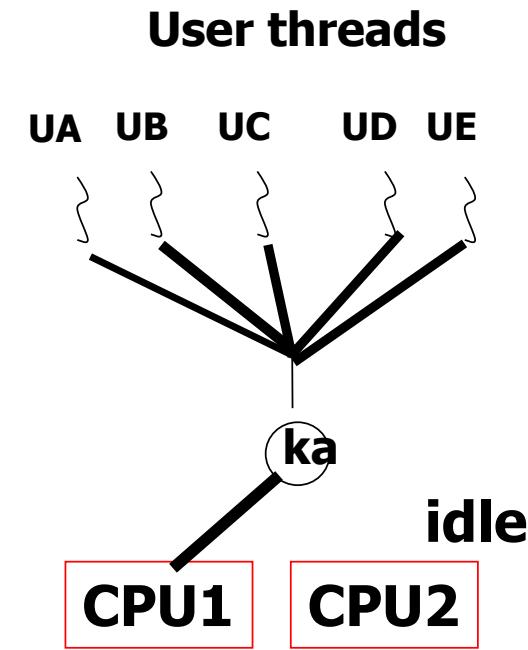
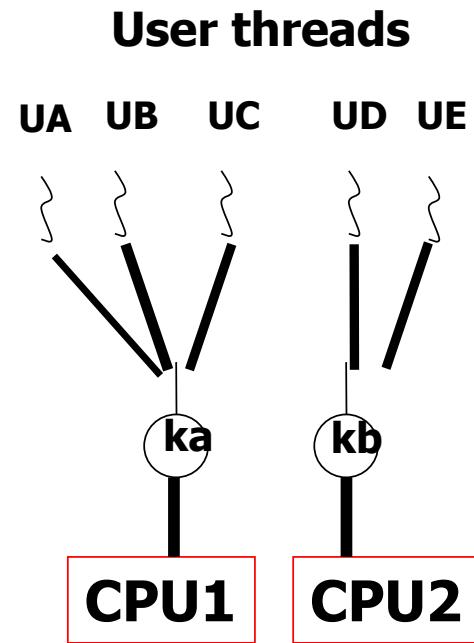
Kernel threads



1. A thread "UA" calls the disk I/O

2. The state of the thread "ka" is changed into the "waiting" state





One kernel thread can be allocated to one processor

