

# Multithreading model

- Many-to-One
- One-to-One
- Many-to-Many

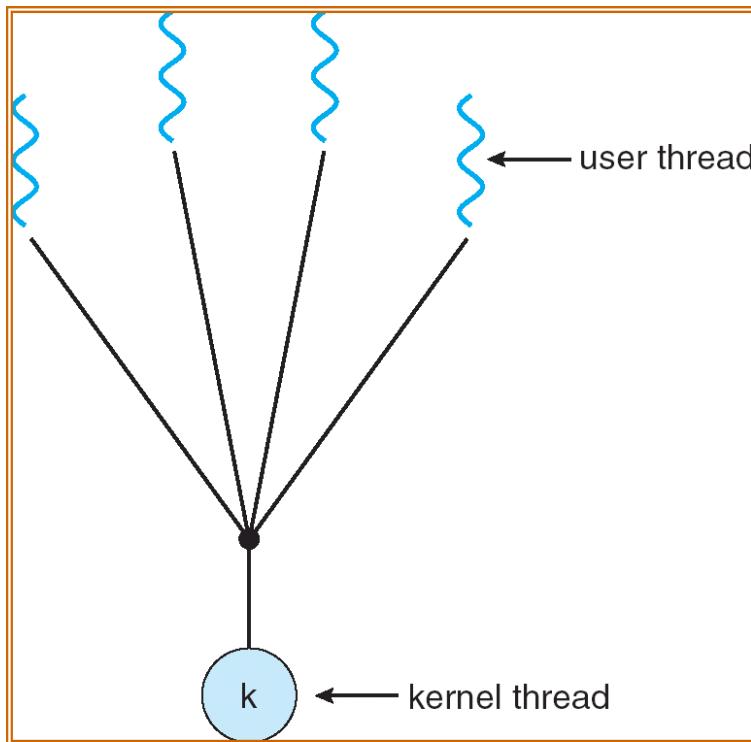
**1. A scheduler is not aware of the existence of user threads**

**2. If a kernel thread is blocked, then its associated user threads are blocked as well.**

**3. One kernel thread can be allocated to one processor**



- Many-to-one model

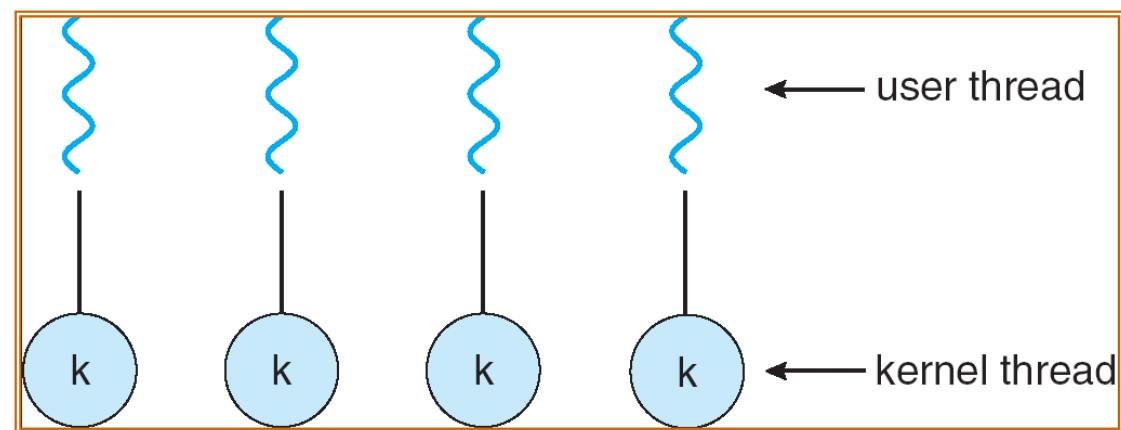


**Very fast, small overhead**

**If one thread calls blocking system call, then the entire process will block**

**Multiple threads are unable to run in parallel on multiprocessors**

## ■ One-to-one model

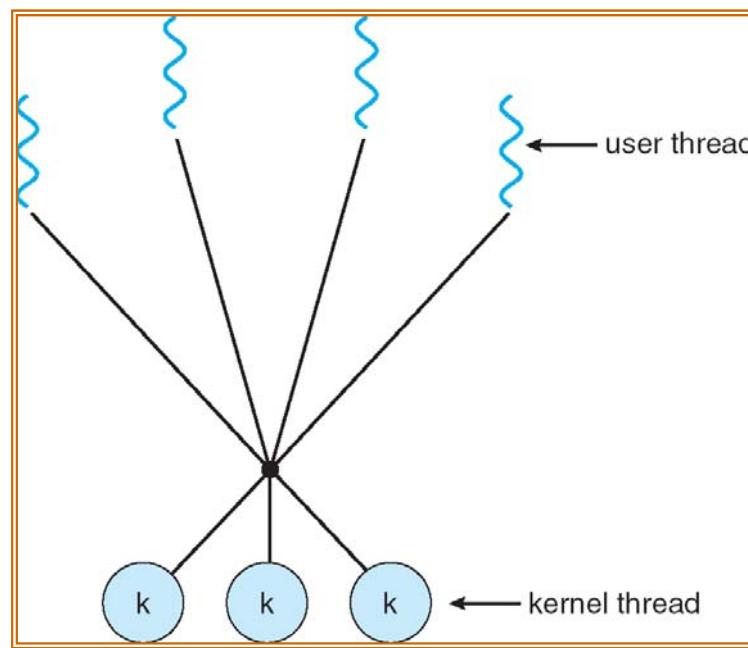


1. Maps each user thread to a kernel thread
2. Another thread can run when a thread makes a blocking system call
3. It allows multiple threads to run in parallel on multiprocessors
4. But creating kernel threads are resource-intensive.



## ■ Many-to-many model

- Multiplexes many user threads to smaller or equal-number of kernel threads





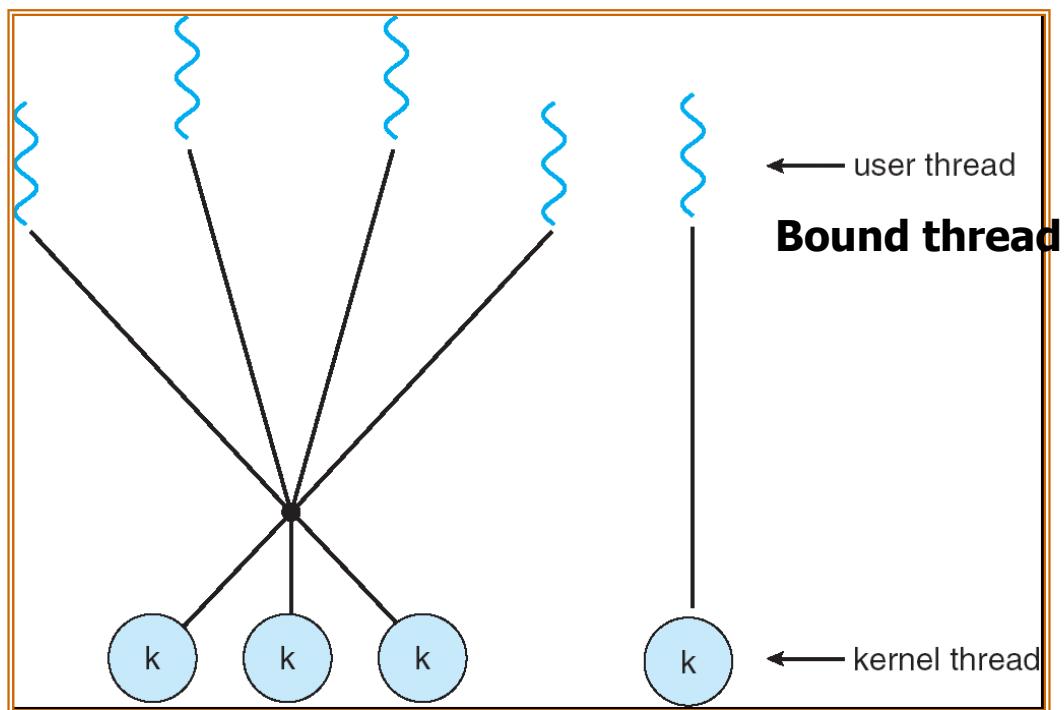
## ■ Many-to-many model

- Hybrid approach between one-to-one model and many-to-one model
- Advantages

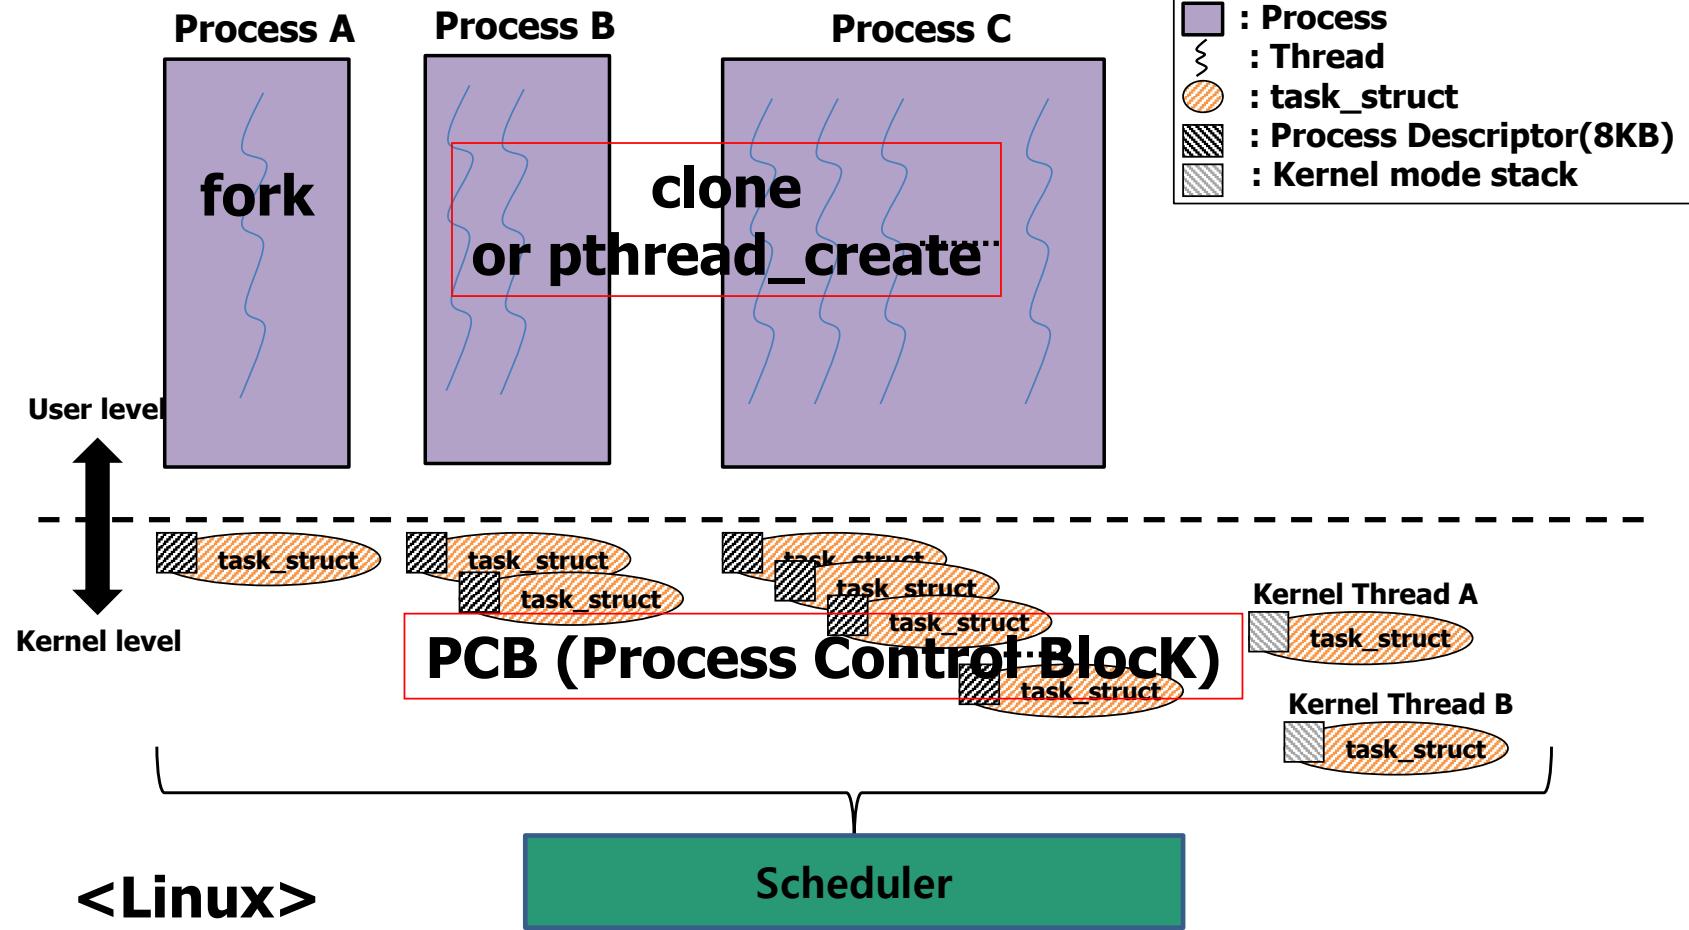
- 1. Developers can create as many user threads as necessary
- 2. The corresponding kernel threads can run in parallel on multiprocessors
- 3. When a thread calls a blocking system call, the kernel can schedule another thread

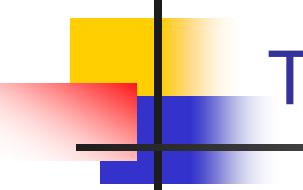
## <Scheduler Activation>

- Two-level model
  - Many-to-many model + one-to-one model



## Linux and windows use one-to-one model



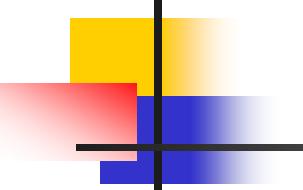


# Thread libraries

- Thread library

- Provides the programmer an API for creating and managing threads
- Two methods;

- To provide a library entirely in user space
  - Invoking a function in the library results in a local function call in user space and not a system call
- To implement a kernel-level library supported directly by the operating system
  - Invoking a function in the library results in a system call

- 
- Three primary thread libraries:
    - POSIX Pthreads
      - May be provided as either a kernel- or user-level library
    - Win32 threads
      - Provided as a kernel-level library
    - Java threads
      - Implemented using a thread library available on the host system
        - For example, on Windows systems, Java threads use Win32 API

## ■ Thread programming example (pthread library)

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */

void *runner(void *param); /* the thread */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of attributes for the thread */

    /* get the default attributes */
    pthread_attr_init(&attr);

    /* create the thread */
    pthread_create(&tid,&attr,runner,argv[1]);

    /* now wait for the thread to exit */
    pthread_join(tid,NULL);

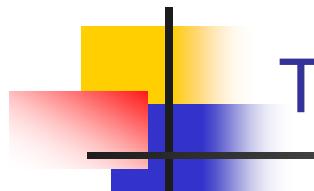
    printf("sum = %d\n",sum);
}
```

### Thread function

```
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    if (upper > 0) {
        for (i = 1; i <= upper; i++)
            sum += i;
    }

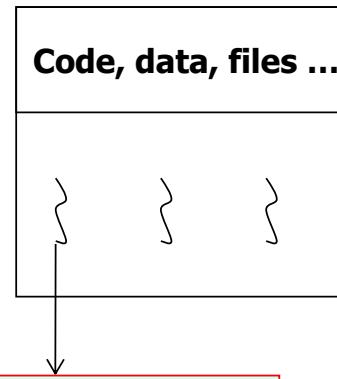
    pthread_exit(0);
}
```



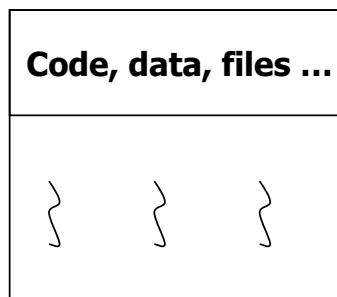
## Threading issues

- Semantics of fork() and exec()
  - Does **fork()** duplicate 1) only the calling thread or 2) all threads?
    - If exec() is immediately called,
      - Only the calling thread is duplicated
    - If exec() is not called,
      - all threads are duplicated

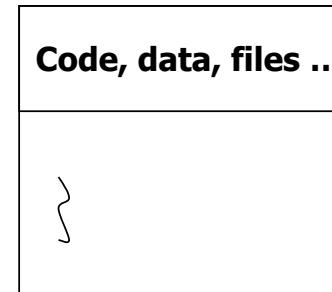




**Fork system call is called**



**If exec system call is not called**



**If exec system call is called immediately**



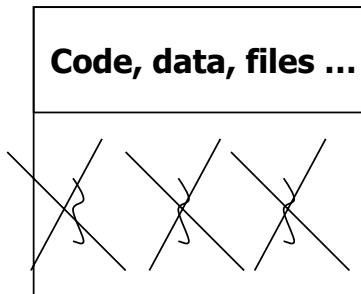


## ■ Thread cancellation

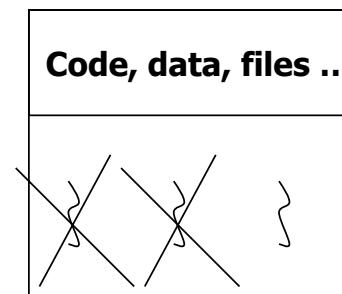
- **Terminating a thread before it has finished**
  - For example,
    - Often, a web page is loaded using several threads
    - When a user presses a button on a web browser that stops a web page from loading any further.
- Target thread
  - A thread that is to be cancelled
- Two general approaches:
  - **Asynchronous cancellation** terminates the target thread immediately
  - **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
    - Safe cancellation points



## Asynchronous cancellation

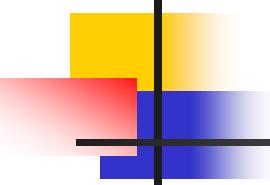


## Deferred cancellation



Check !  
Check !





## ■ Signal handling

- Signals are used in UNIX systems to notify a process that a particular event has occurred;

- Synchronous signal

- Examples: Illegal memory access, divide by zero

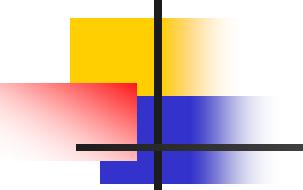
- Asynchronous signal

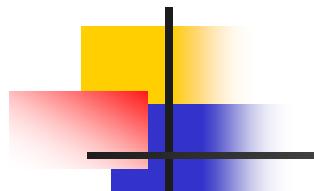
- By an event external to a running process

- signal handler

- A default signal handler

- A user-defined signal handler

- 
- Where should a signal be delivered ?
    - 1. Deliver the signal to the thread to which the signal applies
    - 2. Deliver the signal to every thread in the process
    - 3. Deliver the signal to certain threads in the process
    - 4. Assign a specific thread to receive all signals for the process
  - Examples
    - Division by zero
      - 1.
    - <control><c>
      - 2.
  - POSIX API
    - `pthread_kill (pthread_t tid, int signal)`



**A signal is caused by the division by zero.**

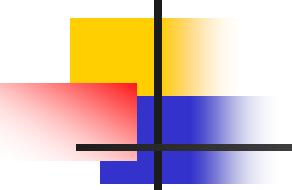
Code, data, files ...



**If the <control><c> button has been pushed**

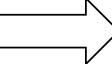
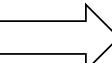
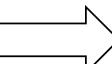
Code, data, files ...

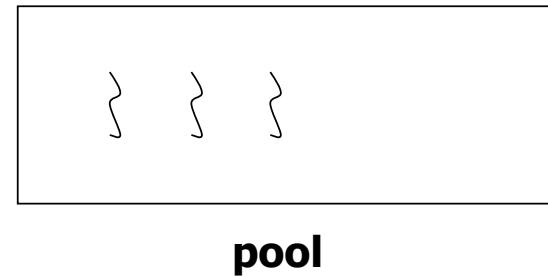


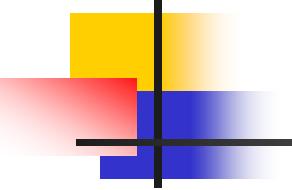


## ■ Thread pools

- Motivation:
  - Unlimited threads may exhaust system resources
- Thread pool
  - Only limited number of threads are admitted to the system
  - Create a number of threads in a pool where they wait for a work
- Advantages:
  - Usually slightly faster to service a request with an existing thread than create a new thread
  - Allows the number of threads in the application(s) to be bound to the size of the pool

- A request** 
- B request** 
- C request** 
- D request**  **Cannot be serviced or wait**

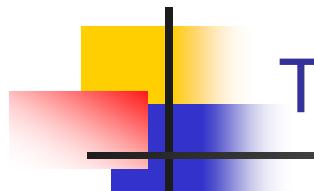




## ■ Thread-specific data

- Allows each thread to have its own copy of data
- Most thread libraries provide some form of support for thread-specific data

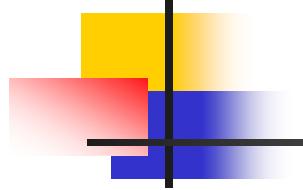




## Thread programming API

- Thread programming API
  - `pthread_create`
  - `pthread_join`
  - `pthread_exit`

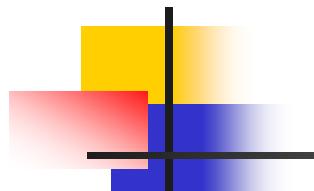




## ■ POSIX thread programming (1)

- Thread creation
  - 1> Prototype
    - #include <pthread.h>
    - int pthread\_create( pthread\_t\* tid, pthread\_attr\_t \*attr, (void \*) f, void \*arg);
  - Roles
    - Creates a new thread and **runs the thread routine f with an input argument of arg**
    - When pthread\_create returns, argument tid contains the ID of the newly created thread





## ■ POSIX thread programming (2)

- Reaping terminated threads
  - Prototype
    - `#include <pthread.h>`
    - `int pthread_join(pthread_t tid, void *thread_return);`
  - Roles
    - `pthread_join` function blocks until thread `tid` terminates
    - It is similar to `wait` function but can only wait for a specific thread to terminate





## ■ POSIX thread programming (3)

- Terminating the threads

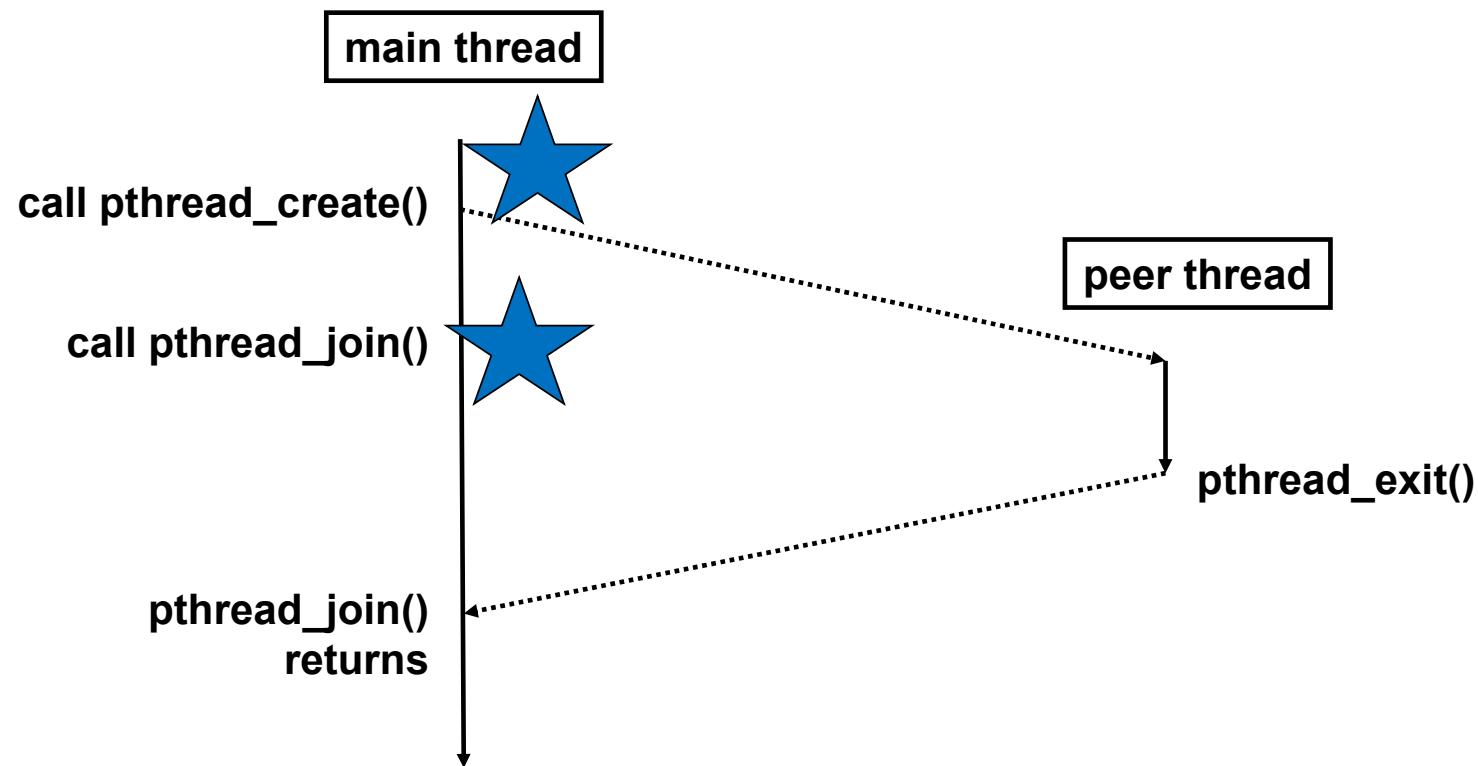
- Prototype

- #include <pthread.h>
    - int pthread\_exit(void \*thread\_return);

- Roles

- Terminating the thread with a return value of `thread_return` that will be transferred to `pthread_join`

## S/W architecture



```

#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int sum; /* this data is shared by the thread(s) */

void *runner(void *param); /* the thread */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */

    /* create the thread */
    pthread_create(&tid,NULL,runner,argv[1]);

    /* now wait for the thread to exit */
    pthread_join(tid,NULL);

    printf("sum = %d\n",sum);
}

```

## Thread function

```

void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    if (upper > 0) {
        for (i = 1; i <= upper; i++)
            sum += i;
    }

    pthread_exit(0);
}

```

